

# Pulling Visitors

Inverting Visitor-Based Control Flow



# Agenda

Daniel J H, works for Mapbox on Graphs

Boost.Graph introduction, from visitors to iterators

<https://github.com/daniel-j-h/cppnow2016> ([git.io/cppnow2016-bgl](http://git.io/cppnow2016-bgl))



# Boost.Graph's Generic Building-Blocks

Data structures (graph types)

Iterators (edges, vertices)

Properties (internal, external)

Algorithms (breadth-first search, dijkstra)

Visitors (examine\_vertex)

# Graph Concepts

Graph types are models for Graph Concepts, determines functionality

IncidenceGraph (source, target, out\_edges)

BidirectionalGraph (in\_edges)

VertexListGraph (vertices)

# Graph Representations

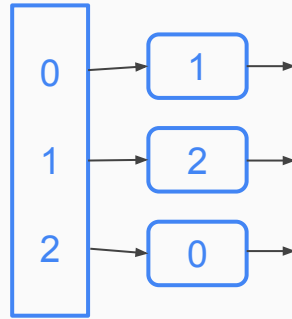
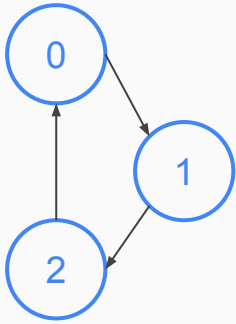
Customizable through template tags: Directed, Undirected, Properties

`adjacency_list`

`adjacency_matrix`

`compressed_sparse_row_graph`

# adjacency\_list<vecS, listS, directedS>



```
using graph_t = adjacency_list<vecS, vecS, directedS>;
```

```
graph_t graph(3);
```

```
add_edge(0, 1, graph);
```

```
add_edge(1, 2, graph);
```

```
add_edge(2, 0, graph);
```



```
struct edge_data_t { int duration = 0; };

using graph_t = adjacency_list<vecS, vecS, directedS, no_property, edge_data_t>;

graph_t graph(3);

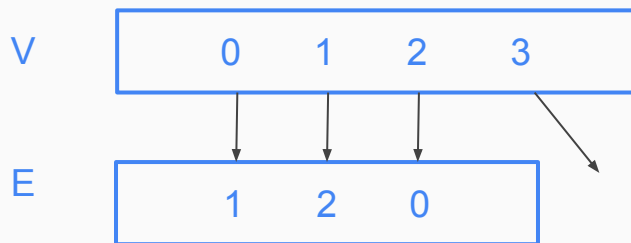
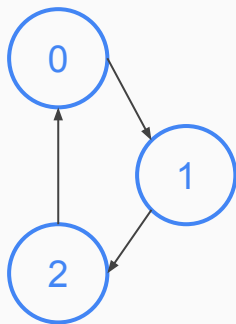
add_edge(0, 1, edge_data_t{100}, graph);


auto duration = [&graph](auto edge) { return graph[edge].duration; };

auto positive = [](auto duration) { return duration > 0; };


auto d = accumulate(edges(graph) | transformed(duration) | filtered(positive), 0);
```

# compressed\_sparse\_row\_graph<directedS>



```
edges(v):  
    first = E[V[v]];  
    last  = E[V[v + 1]];
```

```
using graph_t = compressed_sparse_row_graph<directedS>;
```

```
using vertex_t = graph_traits<graph_t>::vertex_descriptor;
```

```
vector<vertex_t> sources{0, 1, 2};
```

```
vector<vertex_t> targets{1, 2, 0};
```

```
auto tag = construct_inplace_from_sources_and_targets;
```

```
graph_t graph{tag, sources, targets, 3};
```

# Graph Algorithms

Visitors provide algorithm customization points

Graph Walk (bfs, dfs)

Shortest Paths (dijkstra, a-star)

Max-Flow / Min-Cut (edmonds\_karp\_max\_flow)

```
struct discover_visitor : default_bfs_visitor {  
    void discover_vertex(const vertex_t vertex, const graph_t&) {  
        cout << vertex << endl;  
    }  
};
```

```
vertex_t source{0};
```

```
breadth_first_search(graph, source, visitor(discover_visitor{}));
```

# Use-Case Bidirectional Dijkstra

Baseline router to compare against

Start first search on graph from source

Start second search on reversed graph from target

Step both searches (ping-pong) until they meet in the middle

# Problem: how to stop and resume visitors

```
vertex_t middle;
```

```
async(dijkstra_shortest_path(graph, source, visitor(ping_pong{middle}));
```

```
async(dijkstra_shortest_path(rev_graph, target, visitor(ping_pong{middle}));
```

# Coroutines for Cooperative Multitasking

Bind coroutine to visitor, get lazy Dijkstra generator for free

No explicit synchronization, no threads (concurrency != parallelism)

Aha Moment: can be stopped, can be resumed, proper iterators (stdlib)

Technique works for all visitors, and especially well for Boost.Graph



```
using coro_t = coroutines::asymmetric_coroutine<vertex_t>;
```

```
struct dijkstra_stepwise : default_dijkstra_visitor {
```

```
    dijkstra_stepwise(coro_t::push_type& sink_) : sink(sink_) {}
```

```
    void examine_vertex(const vertex_t vertex, const graph_t&) const {
```

```
        sink(vertex);
```

```
    }
```

```
    coro_t::push_type& sink;
```

```
};
```

```
coro_t::pull_type lazy_forward_vertices{[&](auto& sink) {  
    dijkstra_shortest_paths_no_color_map(graph, source,  
        weight_map(get(&edge_data_t::distance, graph))  
        .predecessor_map(forward_prev_map)  
        .visitor(dijkstra_stepwise{sink}));  
}};
```

```
while (lazy_forward_vertices && lazy_backward_vertices)  
    // lazy_forward_vertices.get(); lazy_forward_vertices();
```

```
auto poi = [&graph](auto vertex) { return has_poi(vertex, graph); };
```

```
auto it = find_if(lazy_forward_vertices, poi);
```

```
if (it != end(lazy_forward_vertices))
```

```
    std::cout << *it << std::endl;
```

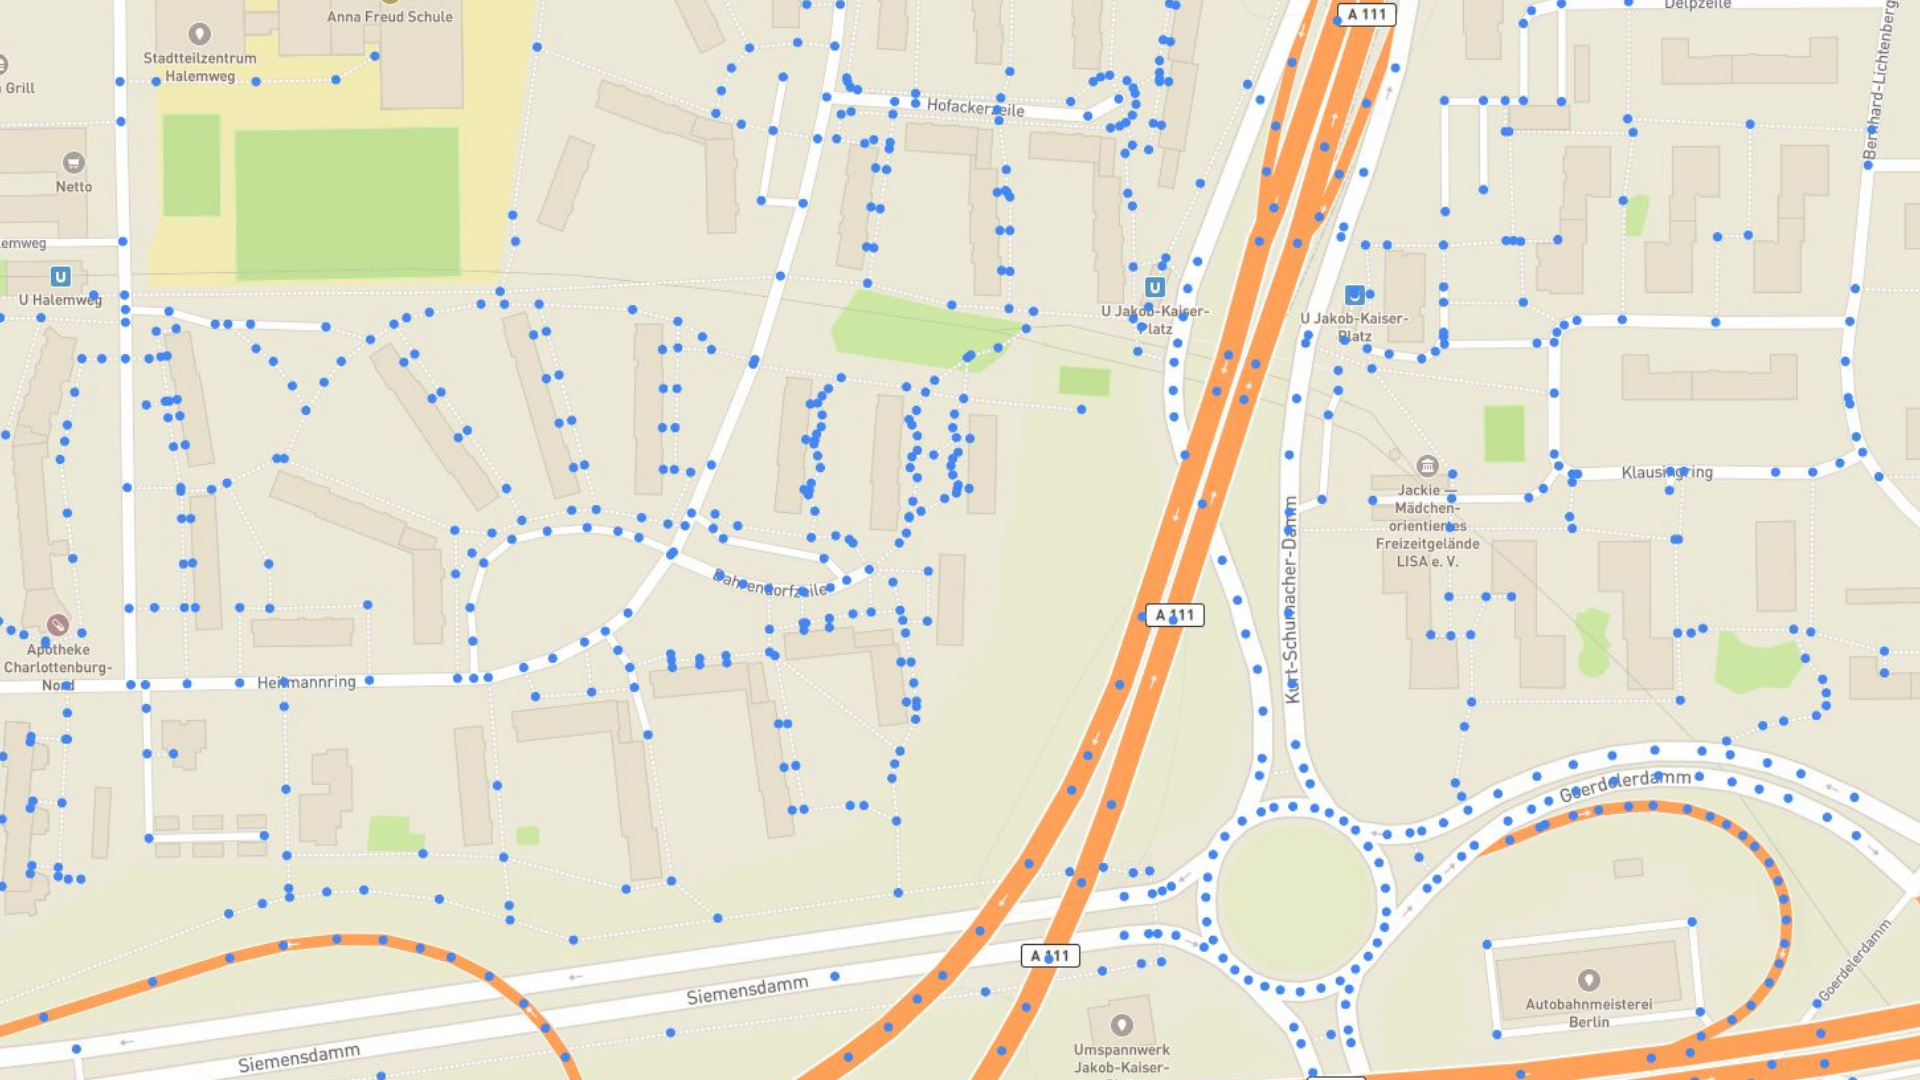
# Give Boost.Graph a try!

OpenStreetMap (extract nodes and ways: libosmium)

Construct graph, add properties (location, Boost.Geometry Haversine distance)

Route on the graph (Boost.Geometry's RTree for initial coordinate lookup)

Visualize the graph (simplification: tippecanoe)



# Our Take Aways

Powerful trio: Boost.Graph + Boost.Range + Boost.Geometry

Switching to 32 bit vertex and edge index types in CSR (size\_t default)

Parallel Boost.Graph vs. r3.4xlarge (120 GB RAM), r3.8xlarge (250 GB RAM)

**Boost.Graph + Boost.Coroutine: from visitors to generators, stdlib integration**